

Bayesian Approaches to Inverse Problems in Astrophysics and Cosmology

lecture 3

Dr. Prashin Jethwa
Institut für Astrophysik

Lecture 3

- Deep-dive into the “gold-standard” inference algorithm: HMC/NUTS
- Gaussian Processes (GPs)

By the end of this lecture you should be able to:


- Describe what the HMC/NUTS algorithm does
 - specifically with reference to: step-size, target_acceptance_probability, warm-up stage, divergent transitions
- Explain the convergence diagnostics: N_{eff} , R_{hat} , $N_{\text{divergences}}$
- Suggest strategies to fix HMC/NUTS inference with poor diagnostics
- Understand what it means that GPs are priors over function space
- Use GPs for inference with the JAX-based `tinygp` package

Markov-Chain Monte Carlo Sampling


- Reminder from last lecture:
- There are several different types of MCMC algorithm
 - with a different transition kernel $T(\boldsymbol{\theta}_m, \boldsymbol{\theta}_{m+1})$
 - [online demonstration](#)
- Compare two:
 - Random Walk Metropolis Hastings
 - HMC/NUTS: Hamiltonian Monte Carlo + No-U-Turn Sampler

Hamiltonian Monte Carlo (HMC) & No U-Turn Sampler (NUTS)



The gold-standard inference algorithm




Target distribution: $p(\mathbf{x})$



Target distribution: $p(\mathbf{x})$




Goal: construct a MCMC transition operator that will draw many uncorrelated samples from the typical set



Random Walk Metropolis Hastings

proposes a transition to a Gaussian blob centered around current position.
This becomes very inefficient for high dimensions


Random Walk Metropolis-Hastings
does not scale well to high dimensions.



Solution: use a transition which moves *through* the typical set towards new, unexplored neighborhoods.


How to transition along the typical set

- Say we have target distribution $p(\mathbf{x})$
- The gradient of the target distribution $\nabla p(\mathbf{x})$ defines a vector field
- What would happen to a particle which follows a trajectory along the gradient?




How to transition along the typical set


- Say we have target distribution $p(\mathbf{x})$
- The gradient of the target distribution $\nabla p(\mathbf{x})$ defines a vector field
- What would happen to a particle which follows a trajectory along the gradient?
 - It will end up at the mode
- ... to instead move *along* the typical set, we must give the particles some *momentum*




Momentum: physical analogy with a gravitational potential



Too little momentum
→ particle crashes to the mode



Correct momentum
→ particle orbits along the typical set



Too much momentum
→ particle escapes

The *correct* momentum via Hamiltonian Dynamics

- Introduce extra **momentum variables** $x \rightarrow (x, v)$
- Define the joint distribution over position and momentum: $p(x, v) = p(v | x) p(x)$
- Write this in terms of a Hamiltonian function H $= \exp [-H(x, v)]$

$$\begin{aligned} \rightarrow \quad H(x, v) &= -\log p(v | x) - \log p(x) \\ &:= K(v, x) + V(x) \\ \text{energy} &:= \text{kinetic energy} + \text{potential energy} \end{aligned}$$

- If we take any initial point (x_0, v_0) and evolve it via **Hamilton's equations**:


$$\begin{aligned} dx/dt &= + \partial H / \partial v \\ dv/dt &= - \partial H / \partial x \end{aligned}$$

then the resulting trajectory naturally orbits along the typical set of the target distribution

HMC: the algorithm

1. Start at initial point

x_n




HMC: the algorithm

1. Start at initial point

$$x_n$$

2. Sample a momentum

$$v_n \sim p(v \mid x_n)$$



HMC: the algorithm

1. Start at initial point


$$x_n$$

2. Sample a momentum

$$v_n \sim p(v \mid x_n)$$


3. Use Hamilton's equations to numerically an orbit

$$(x_n, v_n) \rightarrow (x_{n+1}, v_{n+1})$$



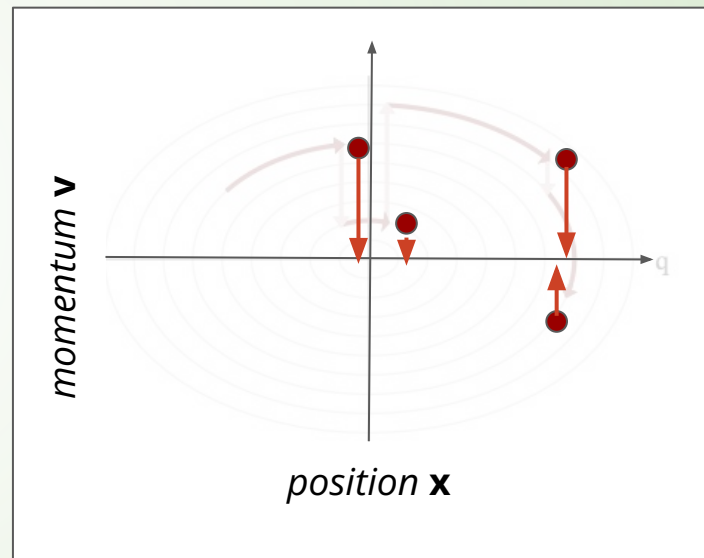
HMC: the algorithm

1. Start at initial point x_n
2. Sample a momentum $v_n \sim p(v \mid x_n)$
3. Use Hamilton's equations to numerically an orbit $(x_n, v_n) \rightarrow (x_{n+1}, v_{n+1})$
4. Accept/reject the new point...
5. Repeat for desired number of samples



HMC: the algorithm

1. Start at initial point x_n
2. Sample a momentum $v_n \sim p(v \mid x_n)$
3. Use Hamilton's equations to numerically an orbit $(x_n, v_n) \rightarrow (x_{n+1}, v_{n+1})$
4. Accept/reject the new point...
5. Repeat for desired number of samples
6. Project samples down to position space i.e. forget the momenta $(x, v) \rightarrow x$



What is $p(v | x)$?

- The choice of $p(v | x)$ - i.e. the kinetic energy term in the Hamiltonian - is a free choice - different possibilities exist
- Most common choice: Euclidean-Gaussian Kinetic Energy:

$$p(v | x) = N(v | 0, M)$$

i.e. the kinetic energy

$$\begin{aligned} K(v, x) &= -\log p(v | x) \\ &= \frac{1}{2} v^T M v + \log |M| \end{aligned}$$

where M is the **mass-matrix**, i.e. tuning-parameters of the algorithm

- M is tuned to re-scale/rotate the parameter space distribution into a more standardised frame

Numerically integrating the orbit?

- Take an initial point (x_0, v_0) and evolve it via Hamilton's equations:

$$dx/dt = + \partial H / \partial v$$

$$dv/dt = - \partial H / \partial x := a$$


- When discretized, with some **step-size** ϵ these can be integrated numerically with a leapfrog (AKA kick-drift-kick) integrator:

$v_{i+1/2}$	$= v_i$	$+$	a_i	$\epsilon/2$	"kick"
x_{i+1}	$= x_i$	$+$	$v_{i+1/2}$	ϵ	"drift"
v_{i+1}	$= v_{i+1/2}$	$+$	a_{i+1}	$\epsilon/2$	"kick"


- This integrator is symplectic - i.e. it preserves energy - in contrast to other schemes e.g. Runge-Kutta methods

Setting the step-size ϵ

- To efficiently make large trajectories, we want a large step-size ϵ
- But if ϵ is too large, energy is no longer conserved
 - the trajectory diverges from the true one




Numerical orbit stays near truth




Numerical orbit diverges from truth

Setting the step-size ϵ

- To efficiently make large trajectories, we want a large step-size ϵ
- But if ϵ is too large, energy is no longer conserved
 - the trajectory diverges from the true one
- We can identify these cases by looking at **energy drift** between the trajectory endpoints:
 - $\Delta E = H(x_0, v_0) - H(x_n, v_n)$
 - If $\Delta E > \text{some threshold}$, the trajectory is **divergent**
- Gives an extra step in HMC algorithm...




Numerical orbit stays near truth



Numerical orbit diverges from truth

HMC: the algorithm

1. Start at initial point x_n
2. Sample a momentum $v_n \sim p(v \mid x_n)$
3. Use Hamilton's equations to numerically an orbit $(x_n, v_n) \rightarrow (x_{n+1}, v_{n+1})$
4. Accept/reject new point based on energy drift in trajectory:
 - a. Large drift \rightarrow less likely to be accepted
 - b. Implemented in a way which satisfies the detailed-balance equation
5. Repeat for desired number of samples
6. Project samples down to position space i.e. forget the momenta $(x, v) \rightarrow x$




No U-Turn Sampler (NUTS)

- As well as a step-size, we need to know *how many* steps to take

No U-Turn Sampler (NUTS)

- As well as a step-size, we need to know *how many* steps to take
- We want enough steps to take us away from our initial point




No U-Turn Sampler (NUTS)

- As well as a step-size, we need to know *how many* steps to take
- We want enough steps to take us away from our initial point
- But not so many that we return there



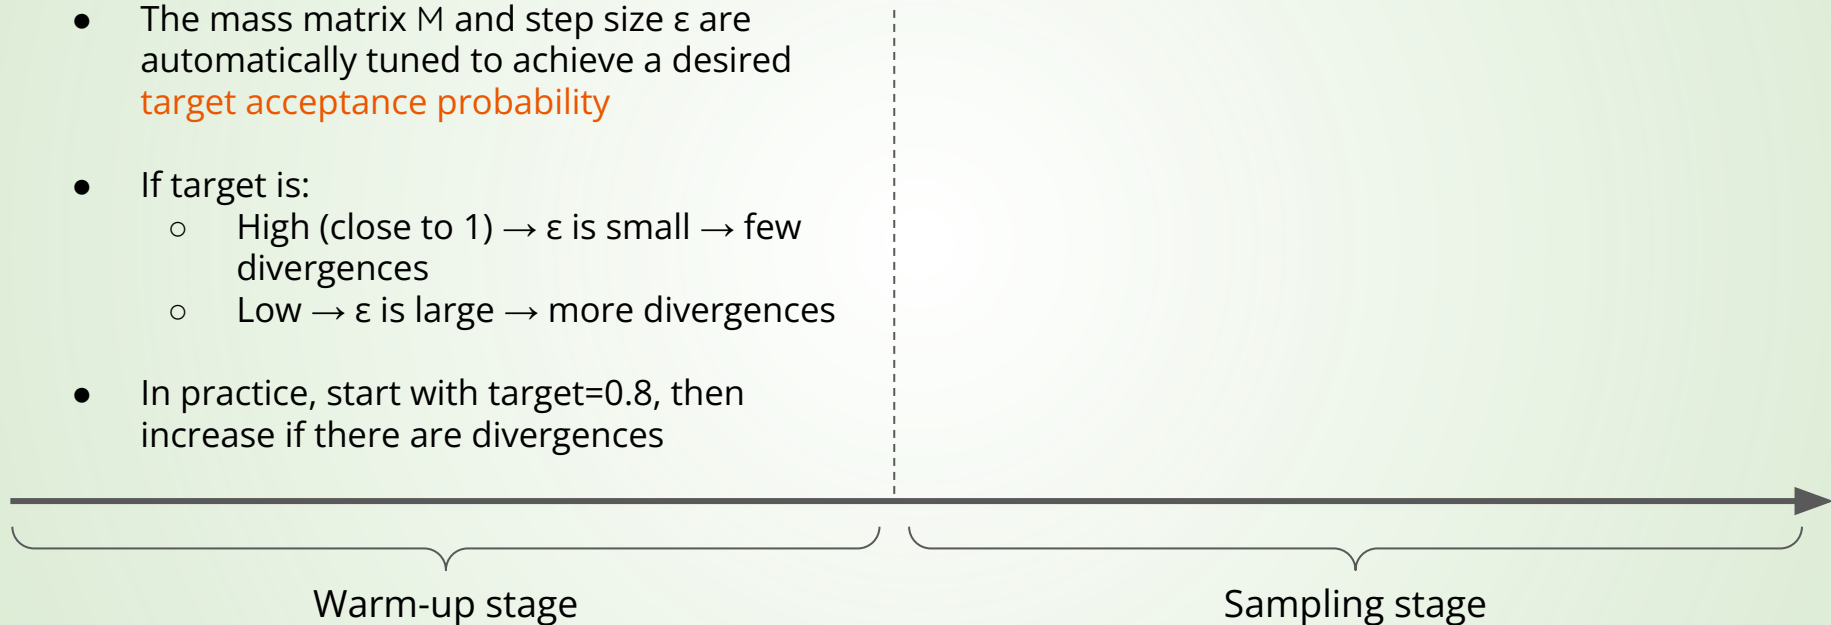
No U-Turn Sampler (NUTS)

- As well as a step-size, we need to know *how many* steps to take
- We want enough steps to take us away from our initial point
- But not so many that we return there
- Solution:
 - Implement a criteria that the trajectories should not perform a “U-Turn”
 - i.e. once they face the opposite direction, stop
 - NUTS was an important ingredient to make HMC work in practice for a wide-variety of problems



Running HMC/NUTS in practice

- The mass matrix M and step size ϵ are automatically tuned to achieve a desired **target acceptance probability**
- If target is:
 - High (close to 1) $\rightarrow \epsilon$ is small \rightarrow few divergences
 - Low $\rightarrow \epsilon$ is large \rightarrow more divergences
- In practice, start with target=0.8, then increase if there are divergences




Number of Steps in MCMC chain

Running HMC/NUTS in practice

- The mass matrix M and step size ϵ are automatically tuned to achieve a desired **target acceptance probability**
- If target is:
 - High (close to 1) $\rightarrow \epsilon$ is small \rightarrow few divergences
 - Low $\rightarrow \epsilon$ is large \rightarrow more divergences
- In practice, start with target=0.8, then increase if there are divergences

- Sampling with HMC/NUTS
 - Uses the tuned mass matrix M and step size ϵ from warm-up
 - Only samples from this stage are used for inference




Number of Steps in MCMC chain

Running HMC/NUTS in practice

- The mass matrix M and step size ϵ are

- The length of the warm-up stage:

- Longer warm-up allows the chain to explore the posterior geometry more fully
- Too short may lead to a mass matrix M and step size ϵ which is tuned too specifically to one region of the posterior, and does not generalise well to other regions
- Useful strategy if inference fails: increase the length of the warm up stage




Number of Steps in MCMC chain


MCMC diagnostics

- Gelman Rubin statistic AKA *Rhat*
 - Has sampling converged onto a stable distribution?
 - If several chains are run (recommended!)
 - measures how well mixed different chains are
 - If one chain is run:
 - measures if first half of the chain well-mixed with the second
 - Recommended: $Rhat < 1.05$ for all parameters

Poorly mixed chains with $Rhat \gg 1$



Chains not well-mixed with themselves or each-other



Chains are well-mixed internally, but not mixed or each-other

MCMC diagnostics

- Gelman Rubin statistic AKA *Rhat*
 - Has sampling converged onto a stable distribution?
 - If several chains are run (recommended!)
 - measures how well mixed different chains are
 - If one chain is run:
 - measures if first half of the chain well-mixed with the second
 - Recommended: $Rhat < 1.05$ for all parameters
- Effective sample size (ESS)
 - Estimate of the number of *independent* samples i.e. accounting for any autocorrelation in the chain
 - Recommended: $ESS > 100$ for all parameters


MCMC diagnostics

- Gelman Rubin statistic AKA *Rhat*
 - Has sampling converged onto a stable distribution?
 - If several chains are run (recommended!)
 - measures how well mixed different chains are
 - If one chain is run:
 - measures if first half of the chain well-mixed with the second
 - Recommended: $Rhat < 1.05$ for all parameters
- Effective sample size (ESS)
 - Estimate of the number of *independent* samples i.e. accounting for any autocorrelation in the chain
 - Recommended: $ESS > 100$ for all parameters
- Divergences:
 - $N = 0$

An example of problematic posterior geometry: the funnel degeneracy


Reminder: Hierarchical Models

- Reminder of hierarchical models:
 - data y collected in K different *contexts*
 - different contexts have different parameters θ_k
 - *Population parameters* mean μ and scale τ control distribution of θ_k
- Limits:
 - $\tau \rightarrow 0$: no variation allowed between contexts
 - $\tau \rightarrow \text{infty}$: large variations allowed




$$\theta_k \sim \text{normal}(\mu, \tau).$$

Funnel degeneracies in Hierarchical Models




- Variation of τ controls the amount of variability between contexts
- Creates a funnel degeneracy with large curvature
- No single step-size will work well for narrow bottom and the wide top



$$\theta_k \sim \text{normal}(\mu, \tau).$$

Funnel degeneracies in Hierarchical Models



- Variation of τ controls the amount of variability between contexts
- Creates a funnel degeneracy with large curvature
- No single step-size will work well for narrow bottom and the wide top

Solutions:

1. Collect/use stronger prior information on the population scale τ to limit the range of the funnel probed during sampling
2. Try the *non-centered parameterisation* for the per-context parameters θ_k

Notebook:


lec3_funnel_degeneracy

Two Parameterizations of the Normal Distribution

Centered Parameterization

sample directly from the normal distribution of interest

$$\theta_k \sim \text{normal}(\mu, \tau).$$




Non Centered Parameterization

sample η from unit normal then shift by μ and re-scale by τ

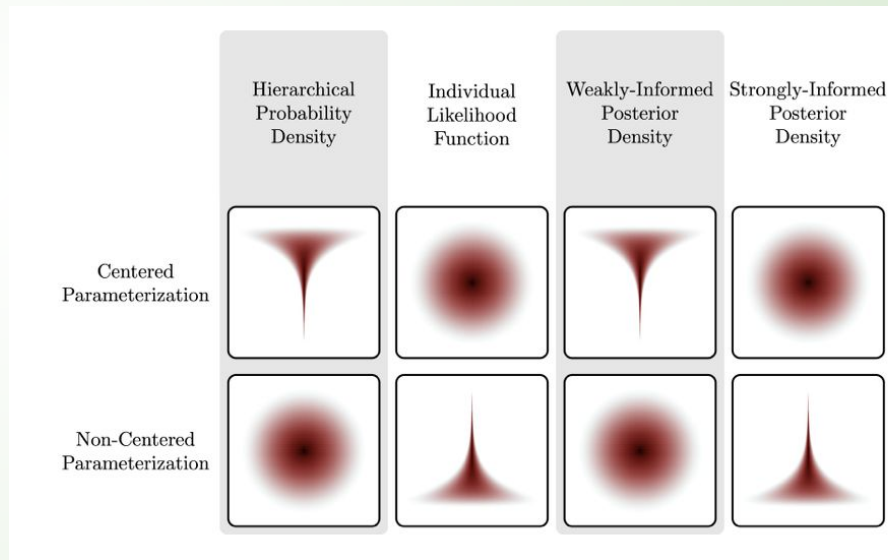
η are normalised deviations of the per-context parameters from the population parameters

$$\eta_k \sim \text{normal}(0, 1)$$
$$\theta_k = \mu + \tau \cdot \eta_k.$$



Why does the centered vs non-centred parameterisation make a difference?

- They represent exactly the same underlying model, so why does a re-parametrisation change behaviour?
- The two parameterisations have posterior geometries which are pathological in complementary regimes
- The regimes depend on the constraining power of the data:
 - τ strongly-informed by data \rightarrow centered best
 - τ weakly-informed by data \rightarrow non-centered best
- Let's walk-through [Michael Betancourt's blogpost](#) for more insight



Why does the centered vs non-centred parameterisation make a difference?

- Centered parameterisation:
 - The prior on context-dependent parameters θ_k and population scale τ displays a funnel
 - If data constrains τ , then funnel pathology is avoided in the posterior
 - When does data constrain τ ?
 - enough individual contexts where θ_k well constrained
 - remaining contexts constrained by partial pooling
- Non-centered:
 - The prior on context-dependent deviation parameters η_k and population scale τ is regular (i.e. no funnel)
 - The *likelihood* now displays a funnel, since $\theta_k = (\eta_k - \mu) / \tau$
 - If data very-constraining, strong funnel is passed from likelihood to posterior
 - If data is weakly constraining, we avoid problematic posterior geometry

Notebook:

lec3_funnel_degeneracy



Practical tips for running HMC/NUTS

- Strategies for dealing with divergent transitions/large Rhat:
 - Use stronger prior information if available
 - Increase the target_acceptance_probability (e.g. from 0.8 to 0.9 to 0.99)
 - Increase the length of the warm-up stage (e.g. from 500 to 1000 to 2000)
 - For hierarchical models displaying a funnel degeneracy, try the non-centered parameterisations
- If effective sample size is low ($ESS < 100$), but other diagnostics look good ($Rhat < 1.05$ and 0 divergences):
 - increase the length of the sampling stage

Gaussian Processes




Motivation

- Linear regression:
 - Two unknown parameters
 - slope m , and intercept c






Motivation

- Linear regression:
 - Two unknown parameters
 - slope m , and intercept c
- What if the data is more complicated than a linear relation?






Motivation

- Linear regression:
 - Two unknown parameters
 - slope m , and intercept c
- What if the data is more complicated than a linear relation?
- Option 1: use a more complex parameterised function
 - e.g. order- N polynomial, with coefficients c_i
 - but we don't know if choice of parameterisation is suitable



Motivation

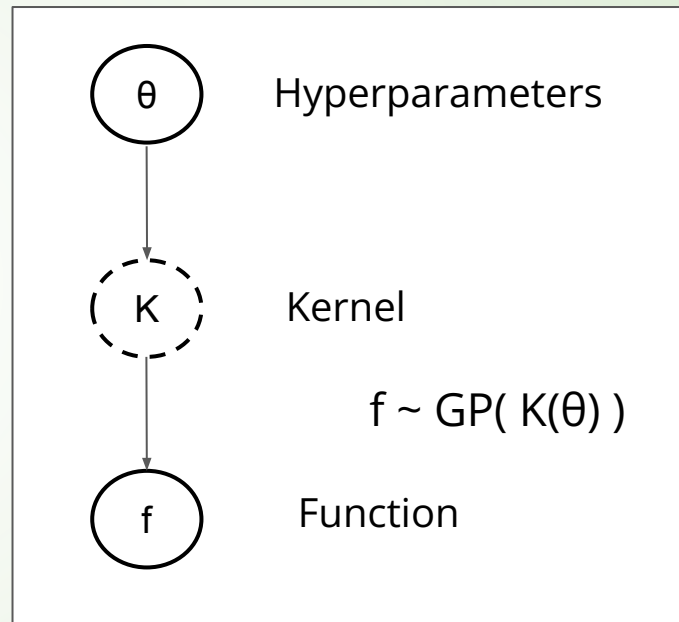
- Linear regression:
 - Two unknown parameters
 - slope m , and intercept c
- What if the data is more complicated than a linear relation?
- Option 1: use a more complex parameterised function
 - e.g. order- N polynomial, with coefficients c_i
 - but we don't know if choice of parameterisation is suitable
- Option 2: directly infer a non-parametric function f which describes the data
i.e. $y = f(x)$



Gaussian Processes (GPs)

- GPs are priors over function spaces
- We can sample *functions* from a Gaussian Process prior
- The GP depends on a kernel K
- The kernel K depends on hyperparameters θ
- Sampling notation:

$$f \sim \text{GP}(K(\theta))$$




The Gaussian Process Kernel

- The kernel K controls the characteristics of the function
- For two points (x_1, x_2) in the input space with separation $dx = |x_1 - x_2|$

$$K(dx)$$


determines how close $f(x_1)$ and $f(x_2)$ are

- i.e. the *smoothness* of f
- Kernel controlled by two hyperparameters:
 - Lengthscale
 - Amplitude



The Gaussian Process Kernel

- **Amplitude:**
 - Larger amplitude \rightarrow the function can have larger vertical variations
- **Lengthscale:**
 - Smaller lengthscale \rightarrow the function can have more high-frequency variations



Notebook:

lec3_gaussian_processes

Gaussian Process: extensions

- We've looked at simple example: 1D regression
- The `tinygp` [documentation](#) contains many more examples:
 - Using more complex kernels e.g. to [model periodic functions](#)
 - Using GPs to model deviations from a [parametric mean function](#)
 - Using GPs within more [complicated numpyro models](#)